

The Social Design Framework

A working session for builders of social platforms. From the social object at the centre, through eight core principles, to the outer layer: Enable, Grow, Protect.



THE CENTRE

Social object · platform & experience
intent



EIGHT PRINCIPLES


Identity · conversations · sharing ·
presence · relationships · reputation ·
groups · agency



THE OUTER LAYER

Enable · grow · protect

One model, three rings, one platform



READING THE RINGS

- **Centre** – the social object that gives people a reason to interact, framed by platform intent (business logic) and experience intent (user value).
- **Inner ring** – eight core design principles. Every social platform makes a choice on each, deliberately or by default.
- **Outer layer** – Enable, Grow, Protect: the holistic read across all principles, with the guiding questions printed on the canvas.

RUNNING THE SESSION

1. Pick one concrete platform – the one you are building. The whole session assesses that product, not platforms in general.
2. Start at the centre: name your social object and both intents before anything else.
3. Work through the eight principles. For each page, discuss the questions and note where your platform stands today.
4. Close with the outer layer – Enable, Grow, Protect – as a holistic read across everything you noted.
5. Dark patterns are what to name and avoid; light patterns are the counter-designs to borrow.

Vocabulary: finitude principles and key concepts

Social media's foundational move was removing edges – no end to the scroll, no boundary between audiences, no natural stopping point. The finitude principles re-introduce them.

THE FINITUDE PRINCIPLES		OTHER CONCEPTS
ExplorativeUI The user moves through a space they can see the shape of, choosing direction – a terrain, not a conveyor belt. Test: can you see where the edges are?	LegibleCuration The user can see, and ideally author, the ranking: why this content, why now, in what order. No black-box recommender.	Social object The shared thing that mediates ties between people – a photo, code, a run, a document (Engeström's object-centred sociality).
EndDesign Genuine endings that release the user – the back page, the last track, "you're all caught up" – not cliffhangers that re-engage.	ConstraintTimeLocking Time budgets enforced as design defaults, authored and held by the user. A platform-imposed lock is paternalism; a user-authored one is sovereignty.	ROI vs. ROX Platform intent is business logic – return on investment. Experience intent is user value – return on experience. The model asks where they diverge.
SpatialUI Bounded contexts (rooms, registers) with their own norms and audiences. Inverts context collapse; grounded in contextual integrity.	DifferentiatedTimeExperiences Distinct tempos – the daily 5-minute check-in, the weekly deep read, the monthly synthesis – instead of one continuous now.	Context collapse Behaviour appropriate to one social context exposed to all at once (boyd/Marwick). Countered by contextual integrity: context-appropriate information flows (Nissenbaum).
SaveForLater Decouples capture from consumption: seeing something interesting and acting on it are separate acts, so asynchronous use is the default.	Completability The interface measures whether you finished your task and left satisfied, not whether you stayed. A graceful exit is a feature.	Bridging vs. bonding capital Bonding reinforces ties within an already-similar group; bridging connects across difference (Putnam). Bonding-only discovery compounds homogeneity.
ReflectiveFriction A deliberate pause at the point of action – before posting, reacting, resharing – restoring the "do I want to do this?" decision.	DefaultMultilingual Genuine multilingualism as baseline expectation, not add-on – generative AI has collapsed the cost to near zero.	Procedural justice Enforcement people accept as legitimate: voice, neutrality, respect, and a trustworthy rationale (Tyler) – not punishment-first moderation.
		Dark & light patterns Designs that extract attention, data, or value from users – and the counter-designs that empower them instead.



Social Object

The shared thing that gives people a reason to interact

04 / 16

The conceptual centre of the framework: the shared thing that mediates ties between people – a photo on Flickr, code on GitHub, a run on Strava, a collaboratively edited document. Assess whether it is rich enough to sustain meaningful sociality, whether users create it or merely consume it, and whether it has value outside the platform.

DARK PATTERNS

- **Object hollowing** – making the social object too thin (likes, status updates) to sustain real sociality
- **Object capture** – locking user-created objects inside the platform with no export
- **Object substitution** – gradually shifting focus from user content to ads/platform content

LIGHT PATTERNS

- **Object depth** – invest the social object with enough structure (provenance, versioning, context) to carry real meaning, not just a status ping
- **Object portability** – full export of user-created objects in open formats, on demand, no friction
- **Object fidelity** – keep platform and ad content structurally separate from, and never crowding out, the object at the center of the experience

QUESTIONS TO CONSIDER

1. What specific thing brings people together on this platform?
2. Is the object something users create or merely consume?
3. Can the social object be exported or preserved if the platform disappears?
4. Does it invite response, remix, or co-creation?

NOTES – WHERE DOES YOUR PLATFORM STAND?

[illegible]



Platform Intent & Experience Intent

Alignment between business logic (ROI) and user value (ROX)

05 / 16

Platform Intent is the business logic (ROI); Experience Intent is user-centred logic (ROX). Revenue tied to engagement is the root cause of the extraction pattern library – infinite scroll, notification harassment, algorithmic amplification – and AI gives that dynamic a more powerful engine. The key test: is the design rooted in finitude, helping users achieve goals and leave satisfied, or in capture?

DARK PATTERNS

- **Misaligned metrics** – optimising for engagement (time spent, clicks) rather than utility or user satisfaction
- **Bait and switch / enshittification** – launching user-centric, then degrading for advertiser or platform revenue
- **Data maximalism** – collecting far more data than the service requires
- **AI-native dark patterns** – hyper-personalised per-user nudges in the model, more dangerous than one-size social-era manipulation

LIGHT PATTERNS

- **Utility metrics** – optimise and report on task completion and satisfaction, not time-on-platform
- **Intent contract** – publish the business model in plain language; changes that break it require re-consent, not silent rollout
- **Data minimalism** – collect only what the service requires; publish a running account of what's collected and why
- **Compleatability metric** – track and publish whether users leave satisfied, not just how long they stayed

QUESTIONS TO CONSIDER

1. Can you explain the business model in one sentence that also describes something genuinely good for users?
2. Are key metrics proxies for user value, or for engagement that proxies for revenue?
3. Who are the customers vs. the users? Whose interests does the design actually serve?
4. If the platform became wildly successful, would its incentives still align with user wellbeing?

NOTES – WHERE DOES YOUR PLATFORM STAND?



Identity

How users reveal, manage, and control who they are

06 / 16

Identity design sets the platform's social contract: real names vs. pseudonymity, context-specific personas, profile granularity, data minimisation. The critical lens is context collapse – behaviour meant for one audience exposed to all – and whether the identity architecture prevents it. Generative AI raises the stakes by inferring and recombining identity signals across contexts at scale.

DARK PATTERNS

- **Forced identity collapse** – requiring a single real-name identity across all contexts, exposing users to context collapse
- **Dark profiling** – collecting identity signals beyond explicit provision (device fingerprinting, shadow profiles, cross-context inference)
- **Identity lock-in** – making identity data non-portable, raising switching costs
- **Compulsory disclosure** – requiring more personal info than the social object requires

LIGHT PATTERNS

- **SpatialUI identities** – support distinct, context-bound personas (work, family, hobby) natively, not as a workaround
- **Profiling transparency** – a visible, plain-language log of every inferred attribute, with one-click deletion
- **Identity portability** – export identity and reputation data in an interoperable format
- **Minimal-disclosure onboarding** – signup asks only what the social object needs; everything else is opt-in later

QUESTIONS TO CONSIDER

1. What is the minimum identity information needed for the social object to function?
2. Can users maintain different context-specific identities (SpatialUI)? Or does the platform enforce a single identity across all social contexts?
3. How does the identity system affect marginalised or vulnerable users?
4. What happens to identity data when a user leaves?

NOTES – WHERE DOES YOUR PLATFORM STAND?



Conversations

How users communicate, respond, and interact

07 / 16

Conversation design shapes whether a platform produces deliberation or provocation, depth or superficiality. It covers what kinds of exchange the architecture enables, whether disagreement escalates or deliberates, notification pressure, and whether there is a deliberate pause before posting or amplifying. AI-generated content at scale makes governance for synthetic posts and personas part of the design.

DARK PATTERNS

- **Rage amplification** – algorithmically surfacing the most emotionally provocative replies to maximise engagement
- **Notification harassment** – aggressive push patterns pulling users back compulsively
- **Frictionless amplification** – one-click reactions and reposts with no ReflectiveFriction before the action
- **AI content laundering** – no governance for AI-generated posts, synthetic personas, or automated engagement

LIGHT PATTERNS

- **ReflectiveFriction on reply** – a brief pause or preview before a heated reply posts
- **Calm notifications** – batched, user-scheduled digests by default; real-time push is opt-in
- **Depth-first threading** – reward formats that support sustained dialogue over one-line dunks
- **AI provenance labelling** – unspoofable labels on AI-generated posts and bots, plus rate limits on synthetic engagement

QUESTIONS TO CONSIDER

1. What kinds of conversation is this platform designed to enable?
2. How does the conversation architecture handle disagreement – does it escalate or deliberate?
3. Does the notification system create pressure for immediate response?
4. Is there ReflectiveFriction before posting, reacting, or resharing?
5. What governance exists for AI-generated content and synthetic engagement?

NOTES – WHERE DOES YOUR PLATFORM STAND?

[illegible]



Sharing

How content flows, spreads, and creates value

08 / 16

Sharing design determines what circulates, who sees it, how it spreads, and who captures the value. Assess default audiences, push vs. pull content flow, friction before resharing, attribution, and whether users can see why content is surfaced to them. Frictionless virality is the extraction pattern; SaveForLater and ReflectiveFriction work against it.

DARK PATTERNS

- **Frictionless virality** – one-click repost with no context, maximising spread regardless of accuracy or intent
- **Algorithmic amplification of extremity** – rewarding emotional provocation for engagement metrics
- **Value extraction** – users create content, platform captures all economic value
- **Context stripping** – removing content from its original social context, enabling misinterpretation

LIGHT PATTERNS

- **Context-preserving share** – resharing keeps the original framing and thread attached
- **Creator-value split** – a direct mechanism (tips, revenue share, licensing) for creators to capture value
- **SaveForLater** – a first-class read-later queue, separate from posting and sharing
- **LegibleCuration** – a "why am I seeing this" explainer on every algorithmically surfaced post

QUESTIONS TO CONSIDER

1. What is the default audience for shared content?
2. Is content flow algorithmic-push or user-pull? Is LegibleCuration in place?
3. Who benefits from sharing – the creator, the community, or the platform?
4. Is there ReflectiveFriction before sharing – was it read, is context preserved, is attribution clear?
5. Is there a SaveForLater path that decouples encountering content from acting on it immediately?

NOTES – WHERE DOES YOUR PLATFORM STAND?



Presence

Whether users can know if others are available and accessible

09 / 16

Presence covers online indicators, typing signals, read receipts, last-seen and location. The ethical question is who benefits from visibility – the user being visible, or the social pressure it creates for response. Healthy presence design supports different temporal modes: the quick check-in, the focused session, the offline period.

DARK PATTERNS

- **Compulsory availability** – presence indicators non-optional, creating social pressure to respond
- **Guilt mechanics** – read receipts and last-seen used to pressure response times
- **Surveillance by design** – broadcasting location/activity without meaningful user control
- **FOMO engineering** – showing what users are missing to drive compulsive checking

LIGHT PATTERNS

- **Presence as opt-in** – online, typing, and read-receipt indicators default off, enabled per-conversation
- **No-guilt read state** – reading a message never signals the sender unless the user chooses to send that signal
- **FOMO-free digest** – “what you missed” is a calm, on-demand summary, not a push-driven loop
- **Low-cost invisibility** – going offline costs nothing socially – no “last seen 3 days ago” shaming

QUESTIONS TO CONSIDER

1. What presence signals does the platform provide, and who do they serve?
2. Can users fully control their own presence visibility with low social cost?
3. Does the platform create an expectation of constant availability?
4. Is the default presence state opt-in or opt-out?
5. Does the presence design support DifferentiatedTimeExperiences – offline periods, focused sessions, slow modes?

NOTES – WHERE DOES YOUR PLATFORM STAND?

[illegible]



Relationships

How users connect, associate, and form social graphs

10 / 16

Relationship design determines the social graph: reciprocal vs. asymmetric connections, how new ties are discovered, the cost of disconnecting, and whether the graph can be exported. The bridging vs. bonding lens asks whether discovery only ever suggests more of the same – compounding homogeneity – or also builds ties across difference.

DARK PATTERNS

- **Asymmetric power** – follow/follower structures producing celebrity hierarchies and parasocial dynamics
- **Contact harvesting** – aggressively importing address books with dark consent patterns
- **Social lock-in** – making the social graph non-portable (high switching costs)
- **Bonding-only design** – discovery and connection mechanisms that only ever reinforce existing homogeneous ties, with no mechanism for bridging across difference

LIGHT PATTERNS

- **Reciprocity by default** – new connection types default to mutual confirmation; one-way "follow" is a separate, deliberate feature
- **Consent-first import** – contact import is opt-in, one contact at a time, never a bulk silent scan
- **Graph portability** – the full connection graph exports in an open format on request
- **Bridging by design** – discovery features that deliberately surface cross-group ties alongside homophilous ones, building bridging capital rather than bonding capital alone

QUESTIONS TO CONSIDER

1. Are connections reciprocal or asymmetric? What power dynamics does this create?
2. Can users categorise connections and control what each group sees?
3. Can the relationship graph be exported?
4. Does the platform produce genuine community or celebrity hierarchies?
5. Does the connection model build bridging capital (cross-group ties), or only bonding capital (reinforcing existing in-groups)?

NOTES – WHERE DOES YOUR PLATFORM STAND?



Groups

How communities and sub-communities form and govern themselves

12 / 16

Groups are where abstract platform potential becomes concrete community. Assess what kinds of groups users can form, what governance tools exist for decision-making and norm-setting, how groups scale, and whether the platform designs for genuine communities of practice or merely audience collections.

DARK PATTERNS

- **Echo chamber amplification** – algorithmically reinforcing group homogeneity
- **Governance void** – providing group creation without governance tools
- **Platform override** – unilaterally changing group rules without community consent
- **Growth-at-all-costs** – pushing groups to grow beyond their healthy size

LIGHT PATTERNS

- **Bridging prompts** – occasional, optional in surfacing of adjacent groups to counter homogeneity
- **Governance starter kit** – every new group gets built-in tools for roles, voting, and norm-setting
- **Community consent for rule changes** – platform policy changes affecting a group require notice and, where feasible, a group-level override
- **Healthy-size nudges** – tools that help large groups split or sub-structure rather than being pushed to grow indefinitely

QUESTIONS TO CONSIDER

1. What governance tools exist within groups?
2. How do groups handle the transition from small community to larger scale?
3. Can groups set and enforce their own norms independently?
4. Does the platform design for communities of practice or audience collections?

NOTES – WHERE DOES YOUR PLATFORM STAND?



Agency

How users take control of their experience and protect their autonomy

13 / 16

Agency is the user's capacity to protect their own attention, wellbeing, and autonomy – the dimension where the finitude principles converge. Assess feed control, algorithmic transparency, data export and account deletion, genuine endings, and user-held time budgets. As platforms integrate AI agents, agency extends to human-in-the-loop checkpoints, consequence-gating, and audit trails.

DARK PATTERNS

- **Illusory control** – settings that appear to give control but are ineffective or quietly undermined
- **Roach motel** – easy to join, deliberately difficult to delete or export
- **Algorithmic opacity** – black-box recommendation with no LegibleCuration
- **No EndDesign** – the platform is engineered for infinite continuation with no genuine endings

LIGHT PATTERNS

- **Real controls** – every visible setting demonstrably changes behaviour; no cosmetic toggles
- **One-step exit** – account deletion and full data export as easy to find and use as sign-up
- **EndDesign** – genuine "you're caught up" states that stop the feed and release the user
- **User-held ConstraintTimeLocking** – time budgets are default-on and user-authored, not a buried opt-in

QUESTIONS TO CONSIDER

1. Can users choose between algorithmic and chronological feeds? Is LegibleCuration in place?
2. How easy is it to leave the platform with your data – is deletion easier or harder than joining?
3. Does the platform have EndDesign – genuine endings, completable sessions?
4. Can users author their own ConstraintTimeLocking? Do they hold the key?
5. If AI agents are present: is there human-in-the-loop oversight, consequence-gating, and an audit trail?

NOTES – WHERE DOES YOUR PLATFORM STAND?



Enable: Foundational Health

Are the foundational conditions for healthy social life established?

14 / 16

Enable asks whether the conditions for healthy social life are established: does the architecture make constructive participation the default, are governance frameworks transparent and participatory, do incentives align with community values? Beneath all dimensions sits one structural question: does the platform re-introduce edges – temporal, contextual, directional – or remove them for capture?

How must the platform's structures and incentives be designed to make constructive participation the easiest and most rewarding choice?

What kind of culture are we intentionally cultivating, and how can rules and governance frameworks reinforce it without stifling it?

DARK PATTERNS

- **Weak governance** – unclear rules and opaque decision-making
- **Misaligned incentives** – rewards that encourage bad behaviour
- **Poor onboarding** – newcomers don't understand norms
- **Fragmented coordination** – concepts work against each other

LIGHT PATTERNS

- **Published governance** – rules, enforcement, and appeals are public and consistently applied
- **Wellbeing-aligned incentives** – internal incentives tied to user-reported satisfaction, not raw engagement
- **Norm-setting onboarding** – the first-run experience actively teaches community norms
- **Coordinated design review** – dimensions are reviewed together for one coherent experience, not shipped in isolation

QUESTIONS TO CONSIDER

1. Do the foundational design choices align with democratic values?
2. Are the platform's rules clear and fairly enforced?
3. Does the incentive structure encourage positive behaviour?
4. Is governance participatory and transparent?

NOTES – WHERE DOES YOUR PLATFORM STAND?



Grow: Sustainable Development

Does the platform create value sustainably without extraction?

15 / 16

Grow addresses how the platform sustains itself and develops over time: whether growth creates meaningful value or extracts attention, whether the business model stays aligned with user wellbeing, and whether network effects and scale benefit everyone – or degrade the experience for those already there.

How can growth remain financially sustainable without compromising the public good?

How does the platform create meaningful value for its participants while contributing positively to society at large?

DARK PATTERNS

- **Extractive growth** – platform captures all value, users capture none
- **Degrading experience** – as platform grows, quality drops for existing users
- **Network lock-in** – switching costs trap users despite poor service
- **Algorithmic enshittification** – quality degrades to push premium features

LIGHT PATTERNS

- **Shared-value growth** – a visible mechanism where platform growth measurably benefits existing users, not just the platform
- **Quality floor** – a committed, monitored quality bar that doesn't erode with scale
- **Low-friction exit** – switching costs kept low; interoperability treated as a trust-building retention strategy, not lock-in
- **Anti-enshittification commitment** – public metrics tracking whether core features degrade to push premium tiers

QUESTIONS TO CONSIDER

1. Does growth create value for users and community, or just for the platform?
2. Is the business model sustainable without becoming exploitative?
3. Do network effects work in users' favour or against them?
4. Does quality improve or degrade as the platform scales?

NOTES – WHERE DOES YOUR PLATFORM STAND?



Protect: Immune System Health

Is the platform's immune system strong, fair, and adaptive?

Protect describes the platform's immune system: how it identifies and responds to threats, internal and external, while balancing safety and openness. Procedural justice is the key lens – enforcement is accepted as legitimate when it offers voice, neutrality, respect, and a trustworthy rationale, rather than punishment alone.

How might scaling, hidden dynamics, or concentrations of power undermine trust or enable manipulation over time?

What structural conditions and incentive systems must be designed to ensure safety, fairness, and long-term resilience?

DARK PATTERNS

- **Security theatre** – visible but ineffective safety measures
- **Disproportionate moderation** – rules applied unfairly to marginalised groups
- **Captured governance** – safety rules designed to benefit the platform, not users
- **Deterrence-only enforcement** – punishment-first moderation with no voice, explanation, or path to appeal

LIGHT PATTERNS

- **Measured safety** – published, audited effectiveness data for safety features, not just their existence
- **Equitable moderation** – outcomes tracked and audited for disparate impact across groups
- **Voice** – a genuine, accessible appeal or reconsideration path before an enforcement decision is treated as final
- **Trustworthy rationale** – a clear, specific explanation of why an action was taken, not a generic policy citation

QUESTIONS TO CONSIDER

1. Are safety mechanisms fair and transparent?
2. Does moderation protect vulnerable users or silence them?
3. Is the community able to contribute to safety?
4. Are the tradeoffs between safety and freedom clearly acknowledged?
5. Does the enforcement system offer voice, neutrality, respect, and a trustworthy rationale (procedural justice), or is it purely punitive?

NOTES – WHERE DOES YOUR PLATFORM STAND?
